

The power of two paths in grid computing networks

Wouter ibens
University of Antwerp

May 10, 2012

Abstract

a

Contents

1	Setup	3
1.1	Forwarding algorithms	4
1.2	Forward to neighbour	4
1.2.1	Forward right	4
1.2.2	Left/Right forward	4
1.2.3	Random Left/Right forward with parameter p	4
1.2.4	Position-dependant forwarding	4
1.3	Forward anywhere	4
1.3.1	Random unvisited	5
1.3.2	Round Robin unvisited	5
1.3.3	Coprime offset	5
1.3.4	Random Coprime offset	5
2	Simulation	5
2.1	Measure	6
2.2	Results	6
3	Numerical Validation	9
3.1	Forward Right	9
3.2	Random Left/Right forward with parameter p	10
3.3	Random Coprime offset	10
3.4	Random Unvisited	10
3.5	Lumped states	10
3.6	Equivalent techniques	11
4	Conclusion	11
A	Simulator source code	11
B	MATLAB Numerical evaluation code	11

Introduction

Write at the end

1 Setup

We are using a ring-structured network of N nodes. Each node is connected to two neighbours, left and right. The purpose of these nodes is to process incoming jobs. When a node is busy while a job arrives, it must forward to another node. When a job has visited all nodes and none of them was found empty, the job is dropped.

Jobs have an arrival time, a length and optional metadata. They arrive at each node independently as a poisson process at rate λ . Their length is exponentially distributed with mean μ (unless otherwise noted, assume $\mu = 1$). Although each job has a length, this length may not be known in advance. Finally, the metadata is optional and may be used by the nodes to pass information among the job (e.g. a list of visited nodes).

Nodes can use different algorithms to determine whereto a job will be forwarded. The performance of these algorithms is the main focus of this thesis. Different techniques will be discussed and simulated. Afterwards, some results of the simulation will be validated. Note that the cost of forwarding a job is neglected. Together with the presumption a job must visit each node before being dropped, this means a job arriving at any node will be processed if and only if at least one server is idle.




Figure 1: A ring structured network

The performance of a forwarding algorithms is measured by the average number of hops a job must visit before being executed. The goal of the algorithms is to minimize this number by spreading the load evenly along the ring.

1.1 Forwarding algorithms

Nodes that must forward a job must choose another node of the ring. Nodes have no information about other nodes, so is has no idea whether the node is idle or busy. The algorithms are grouped in two categories: forward to neighbour and forward anywhere. The first techniques allows a busy node to forward an incoming job to either its left or right neighbour, where the latter may forward these jobs to any node in the ring. Since the amount of dropped jobs is equal for each forwarding algorithm. These jobs will be ignored when computing the average number of hops. One should note that the loss rate of jobs in the system is the same a the Erlang-b loss rate.

1.2 Forward to neighbour

1.2.1 Forward right

A busy node using this technique will forward a job to its right neighbour. The job will keep travelling clockwise until an idle node is found, where it will be processed. This algorithms is used as base line in all further tests.

1.2.2 Left/Right forward

A variant to the previous algorithm is the Left/Right forward technique. Instead of forwarding each job to its right neighbour, a busy node will alternate the direction after forwarding such a job. To avoid a job coming back, this initial direction is saved in the job's metadata. Busy nodes receiving a job from another neighbour must forward it the same direction as specified in the job's metadata.

1.2.3 Random Left/Right forward with parameter p

This technique is a variant of the Left/Right forward algorithm. However, instead of alternating the direction for each new job, a node will forward a job to its right with probability p and to its left with probability $1 - p$. As the previous technique, the direction is saved in the job's metadata and subsequent nodes must maintain this direction when forwarding.

1.2.4 Position-dependant forwarding

As shown in figure 1, each node in the ring has an unique ID. Except the for node with id 0 and $N - 1$, neighbouring are succeeding. When nodes uses this algorithm, nodes will always forward a new job in the same direction: to the right when the node's id is even, to the left otherwise. As previous algorithms, the direction is saved in the job's metadata and this direction must be used if other nodes must forward the job.

1.3 Forward anywhere

The ring strucure can be used in real networks, however in many cases the ring is no more than a virtual overlay over another structure (e.g. the internet). In these networks each node is able to connect to each other node and other forwarding algorithms can be used.

1.3.1 Random unvisited

The Random unvisited algorithm is the most basic algorithm in this category. Everytime a job is forwarded, a list of unvisited nodes is generated and a random node is choosen from this list. The current node is added to the list of visited nodes, which is found in the job's metadata.

1.3.2 Round Robin unvisited

This algorithm is similar to Random unvisited, but instead of choosing the next node at random a different technique is used. When a node must forward a job, it saves its id. When another job is forwarded, it will be forwarded to the the saved id + 1. However, when that id is a visited node, the job will be forwarded to the next node that is unvisited.

1.3.3 Coprime offset

Another algorithm is Coprime offset. This algorithm generates a list of all numbers smaller than N , and coprime to N . The first time a job is forwarded the next number of this list is selected. This is the job's forward offset and saved in the its metadata. When a job is forwarded, it is sent exactly this many hops farther. Because this number and N are coprime, it will visit all nodes exactly once before being returned to its originating node.

Example: Consider a ring size of $N = 10$ in which every node is busy. The list of coprimes is than generated: 1, 3, 7, 9. Assume a job arrives at node 3 and the last time node 3 forwarded a job it was given offet 1. Because this node is busy, the next number on the list (3) is selected and saved in the job's metadata. All nodes are busy so the job visits these nodes before being dropped: 3 (arrival), 6, 9, 2, 5, 8, 1, 4, 7, 0. Node 0 will drop the job because the next node would be 3, which is the node on wich the job arrived.

1.3.4 Random Coprime offset

The Random Coprime offset algorithm is almost equal to Coprime offset. The difference between them is the decision of the offset value. Where it is the next number on the list in Coprime offset, a random value is taken from the list when using Random Coprime offset.

2 Simulation

To evaluate the different algorithms discussed in the previous section, 2 methods will be used. Firstly using a simulation, the second method is the evaluation of this simulation using MATLAB. The validation method is further discussed in section 3.

The simulation is accomplished using a custom simulator. A continuous time simulator is written in C++, using no external requirements but the STL. The source code of the simulator can be found in appenix A or on <http://code.google.com/p/powerofpaths/>.

The simulator can be controlled using a command line interface, its usage is described below.

```

1 Usage: -r -s long -j double -a double -n long -p long -l long -t
   long -h type
2   -r      Random seed
3   -s      Set seed                                (default: 0)
4   -j      Job length                              (default: 1.0)
5   -a      Interarrival time                       (default: 1.0)
6   -n      Ring size                               (default: 100)
7   -p      Print progress interval (default: -1 - disabled)
8   -l      Simulation length                       (default: 3600)
9   -t      Repetition                              (default: 1)
10  -h      Print this help
11  type    right | switch | randswitch | evenswitch | prime |
          randprime | randunvisited | totop

```

Listing 1: Simulator usage description

2.1 Measure

The goal of the algorithms is to distribute the jobs evenly along the ring. This implies the number of hops a job must travel should be low. As a measure for our experiments, we will be using the number of times a job was forwarded before it was executed. Since the number of forwards of a job that could not be executed is the same for each algorithm, and the loss rate of each algorithm is the same, we will not take these jobs into account when computing the average.

It is clear that when the system load approaches 0, the probability that a node is busy will also approach 0 and the average number of forwards will therefore also approach 0. On the other hand, when the load approaches ∞ , each node's probability of being busy will approach 1 and therefore the number of forwards will be $N - 1$ and the job will fail. A system with load > 1 is called an overloaded system.

We will compare each algorithms to a baseline result. The baseline used in this thesis is the Forward right algorithm, meaning that each graph will show its result relative to the Forward right results. The results given by the simulator were obtained using a ring size of 100 and using a random seed for each run.

The absolute performance of the baseline algorithm is shown in figure 2.

2.2 Results

Left/Right Forward

It is intuitively clear that alternating the forwarding direction of arriving jobs should distribute the load better than keeping the same direction. Figure 3 shows the improvement made by the Left/Right forward algorithm over the Forward right method. The performance gain is at least 1% and up to over 4% under medium load.

Random Left/Right forward with parameter p

For $p = 0.5$, one would expect the results of this algorithm being similar to those obtained in the previous simulation. However, it seems the small change in the algorithm worsened the results significantly.

Figure 4 shows the results of this algorithm for $p = 0.5$. How this parameter influenced the performance is shown in figure 5. For $p = 0$, this algorithm is




Figure 2: The Forward Right baseline result




Figure 3: Left/Right

equivalent to the Forward right method. The performance decreases fast when inceasing p , until around 0.4, where is increases a little until arriving at 0.5.

Position-dependant forwarding

This technique groups nodes in virtual clusters. When a job arrives in a node and that node is busy, the job will be forwarded to the other node in the cluster. Jobs leaving a cluster will do this in a random direction ($p = 0.5$). Since the load is concentrated per cluster instead of being distributed over the whole system, this technique performs worse than other techniques The results are represented




Figure 4: Random Left/Right forward with parameter 0.5




Figure 5: Random Left/Right forward with load 0.8

in figure 6.

Random unvisited

This algorithm is the most straight forward and is the best performing from any of these techniques. However, it should be noted that each visited node must be stored into the job's metadata.




Figure 6: Position-dependant forwarding

3 Numerical Validation

To validate the results obtained in the previous section, we modelled the scheduling-techniques into Markov Chains. Using the steady state distribution of these chains, we can derive the average number of hops and the average loss. For N nodes in a ring, the markov chain consists of 2^N number of states, where the bits represent whether a server is busy or idle. To optimize the computation time and memory requirements, we used sparse matrixes for the validation. The validation code is written in MATLAB.

3.1 Forward Right

As for this technique, the matrix representing the Markov chain is easily determined. The example given below is for a ring of 3 nodes. For convenience, the states are representes by their binary form.

$$Q = \begin{matrix} & \begin{matrix} 000 & 001 & 010 & 011 & 100 & 101 & 110 & 111 \end{matrix} \\ \begin{matrix} 000 \\ 001 \\ 010 \\ 011 \\ 100 \\ 101 \\ 110 \\ 111 \end{matrix} & \begin{bmatrix} -3\lambda & \lambda & \lambda & 0 & \lambda & 0 & 0 & 0 \\ \mu & -3\lambda - \mu & 0 & \lambda & 0 & 2\lambda & 0 & 0 \\ \mu & 0 & -3\lambda - \mu & 2\lambda & 0 & 0 & \lambda & 0 \\ 0 & \mu & \mu & -3\lambda - 2\mu & 0 & 0 & 0 & 3\lambda \\ \mu & 0 & 0 & 0 & -3\lambda - \mu & \lambda & 2\lambda & 0 \\ 0 & \mu & 0 & 0 & \mu & -3\lambda - 2\mu & 0 & 3\lambda \\ 0 & 0 & \mu & 0 & \mu & 0 & -3\lambda - 2\mu & 3\lambda \\ 0 & 0 & 0 & \mu & 0 & \mu & \mu & -3\mu \end{bmatrix} \end{matrix}$$

Analogue to the simulation section, this method will be the baseline result in our other results.

3.2 Random Left/Right forward with parameter p

This matrix is very similar to the one above. But we need to take into account the parameters p and $1 - p$ instead of 1 and 0.

$$Q = \begin{matrix} & \begin{matrix} 000 & 001 & 010 & 011 & 100 & 101 & 110 & 111 \end{matrix} \\ \begin{matrix} 000 \\ 001 \\ 010 \\ 011 \\ 100 \\ 101 \\ 110 \\ 111 \end{matrix} & \begin{bmatrix} -3\lambda & \lambda & \lambda & 0 & \lambda & 0 & 0 & 0 \\ \mu & -3\lambda - \mu & 0 & (2-p)\lambda & 0 & (1+p)\lambda & 0 & 0 \\ \mu & 0 & -3\lambda - \mu & (1+p)\lambda & 0 & 0 & (2-p)\lambda & 0 \\ 0 & \mu & \mu & -3\lambda - 2\mu & 0 & 0 & 0 & 3\lambda \\ \mu & 0 & 0 & 0 & -3\lambda - \mu & (2-p)\lambda & (1+p)\lambda & 0 \\ 0 & \mu & 0 & 0 & \mu & -3\lambda - 2\mu & 0 & 3\lambda \\ 0 & 0 & \mu & 0 & \mu & 0 & -3\lambda - 2\mu & 3\lambda \\ 0 & 0 & 0 & \mu & 0 & \mu & \mu & -3\mu \end{bmatrix} \end{matrix}$$

The lumped matrix of Q is equal to the lumped matrix of the example above (3.5), i.e. the matrix defines the exact same problem. However, for $N > 6$ the matrices and so the results of the steady state distribution begin to differ.

3.3 Random Coprime offset

Modelling this technique yields different results for various ring sizes. The performance of this algorithm is very dependant on the number of coprimes that can be used. This technique yields the same results as Forward Right for ring sizes of up 4. For $N = 3$, the matrix Q is identical to Random Left/Right forward with parameter 0.5, as the coprimes of 3 are 1 and 2. Which means forwarding a job left or right, both with the same probability.

3.4 Random Unvisited

This problem can be modelled much more efficiently than the techniques above. Since the next node is chosen randomly, the information we need to save consists only of the number of servers which are currently busy. This problem is analogue to modelling an Erlang-B loss system. The number of states in this Markov Chain is linear to N and is much more dense. For $N = 3$, the matrix is given below.

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \end{matrix} & \begin{bmatrix} -3\lambda & 3\lambda & 0 & 0 \\ \mu & -3\lambda - \mu & 3\lambda & 0 \\ 0 & \mu & -3\lambda - \mu & 3\lambda \\ 0 & 0 & \mu & -\mu \end{bmatrix} \end{matrix}$$

3.5 Lumped states

Except for Random Unvisited, each discribed technique is modelled into a matrix with N^2 states. However, may of these states are redundant: for example, for $N = 3$ the states 001, 010 and 100 all represent one of the nodes being busy. For states where multiple nodes are busy, the space between these servers is

critical information, therefor states are lumped when rotating the bits of one state becomes equal to another state. The states below are analogue and can be lumped:

$$001101 = 011010 = 110100 = 101001 = 010011 = 100110$$

The number of states in a lumped MC is $\frac{1}{N} \sum_{d|N} (2^{N/d} \cdot \text{phi}(d))$ with $\text{phi}(d) = d \cdot \prod_{p|d, p \text{ is prime}} (1 - \frac{1}{p})$ [1]. This result greatly reduces the number of states, however its complexity is still non-polynomial.

The example model in 3.1 can be lumped into the following Markov Chain:

$$Q = \begin{array}{c} \begin{array}{cccc} & 000 & 001 & 011 & 111 \\ \begin{array}{c} 000 \\ 001 \\ 011 \\ 111 \end{array} & \begin{bmatrix} -3\lambda & 3\lambda & 0 & 0 \\ \mu & 3\lambda - \mu & 3\lambda & 0 \\ 0 & 2\mu & 3\lambda - 2\mu & 3\lambda \\ 0 & 0 & 3\mu & -3\mu \end{bmatrix} \end{array}$$

3.6 Equivalent techniques

Lumping states of a Markov Chain produces an equivalent Markov Chain. This can be used to prove some forwarding techniques are equal up to a certain N .

4 Conclusion

Which techniques work best in which environments? Why? Runner up? Why do some techniques don't work as expected?

Acknowledgements

Thanks to everyone

References

- [1] The Online Encyclopedia of Integer Sequences. *A000031*. June 2009. URL: <https://oeis.org/A000031>.

A Simulator source code

dit is de inhoud

B MATLAB Numerical evaluation code

```
1 function [Q] = rightchain(size, rate)
2 %RIGHTCHAIN Generate a Markov Chain that always forwards right
3 %Parameters:
4 %     size    The size of the ring
5 %     rate    The rate of arrivals
```

```

6
7     totalsize = 2^size;
8     Q = sparse(totalsize, totalsize);
9
10    BITS = zeros(1, size);
11
12    for i=1:size
13        BITS(i) = 2^(i-1);
14    end
15
16    for i=0:(totalsize-1)
17        t=0;
18        for b=1:size
19            j=bitxor(i, BITS(b));
20            if bitand(i, BITS(b))
21                Q(i+1, j+1)=1;
22            else
23                r=rate;
24                bt=b+1;
25                while bitand(i, BITS(mod(bt-1, size)+1)) & (bt ~= (
26                    b))
27                    bt=bt+1;
28                    r = r + rate;
29                end
30                Q(i+1, j+1)=r;
31            end
32            t=t + Q(i+1, j+1);
33        end
34        Q(i+1, i+1) = -t;
35    end
36 end

```

Listing 2: rightchain.m

```

1 function [Q] = randswitchchain(size, rate, p)
2 %RANDSWITCHCHAIN Generates a Markov Chain that randomly forward
   left or right
3 %Parameters:
4 %     size      The size of the Markov Chain
5 %     rate      The rate of arrivals
6 %     p          The probability a job is forwarded right
7
8     if nargin < 3
9         p=0.5;
10    end
11
12    totalsize = 2^size;
13    Q = sparse(totalsize, totalsize);
14
15    BITS = zeros(1, size);
16
17    for i=1:size
18        BITS(i) = 2^(i-1);
19    end
20
21    for i=0:(totalsize-1)
22        t=0;
23        for b=1:size
24            j=bitxor(i, BITS(b));
25            if bitand(i, BITS(b))
26                Q(i+1, j+1)=1;

```

```

27         else
28             r=rate;
29             bt=b+1;
30             while bitand(i, BITS(mod(bt-1, size)+1)) & (bt ~= (
31                 b))
32                 bt=bt+1;
33                 r = r + rate*p;
34             end
35             bt=b-1;
36             while bitand(i, BITS(mod(bt-1, size)+1)) & (bt ~= (
37                 b))
38                 bt=bt-1;
39                 r = r + rate*(1-p);
40             end
41             Q(i+1, j+1)=r;
42         end
43         t=t + Q(i+1, j+1);
44     end
45     Q(i+1, i+1) = -t;
46 end

```

Listing 3: randswitchchain.m

```

1 function [ Q ] = rprimechain( size , rate )
2 %RPRIMECHAIN Generate a Markov Chain that chooses a random coprime
3 and uses this as forwarding offset
4 %Parameters:
5 %    size    The size of the ring
6 %    rate    The arrival rate
7
8 totalsize=2^size;
9 rprimes=[];
10
11 for i=1:(size-1)
12     if gcd(size, i) == 1
13         rprimes=[rprimes i];
14     end
15 end
16
17 rpcount = length(rprimes);
18
19 %Q=zeros(totalsize);
20 Q=sparse(totalsize, totalsize);
21
22 for i=0:totalsize-1
23     tot=0;
24     for j=0:size-1
25         k=2^j;
26         if bitand(i,k)
27             Q(i+1, i-k+1) = 1.0;
28             tot=tot+1.0;
29         else
30             c=0;
31             for p=rprimes
32                 current=mod(j-p, size);
33                 while (bitand(i,2^current))
34                     current=mod(current
35                         -p, size);
36                     c=c+1;
37                 end
38             end
39         end
40     end
41 end

```

```

36                                     end
37                                     Q(i+1, i+k+1) = rate + c*rate/
                                         rpcount;
38                                     tot=tot+Q(i+1, i+k+1);
39                                     end
40                                     end
41                                     Q(i+1, i+1) = -tot;
42     end
43 end
44 end

```

Listing 4: rprimechain.m

```

1 function [ Q ] = runvisitedchain( size , rate )
2 %RUNVISITEDCHAIN Generate a Markov Chain that forwards to an
   unvisited node
3 %Parameters:
4 %   size      The size of the ring
5 %   rate      The arrival rate
6
7     rate = rate*size;
8
9     Q = sparse(size+1, size+1);
10
11     Q(1,2) = rate;
12     Q(1,1) = -rate;
13
14     Q(size+1, size) = size;
15     Q(size+1, size+1) = -size;
16
17     for i=2:size
18         Q(i, i-1) = i-1;
19         Q(i, i+1) = rate;
20         Q(i, i) = -Q(i, i+1)-Q(i, i-1);
21     end
22 end
23 end

```

Listing 5: runvisitedchain.m

```

1 function [ avg ] = avghops(Q, d)
2 %AVGHOPS Calculate average number of times a job is forwarded
3 %Parameters:
4 %   Q          The matrix representing a markov chain
5 %Optional:
6 %   d          Debug mode, default=1, disable debug output=0
7
8     if nargin < 2
9         d=1;
10    end
11
12    steady=full(ctmcsteadystate(Q));
13
14    len=length(Q);
15    states=log2(len);
16    avg=0;
17    total=0;
18
19    for i=0:(states-1)
20        c=0;
21        prefix=((2^i)-1) * 2^(states-i);
22        for j=0:(2^(states-i-1))-1

```

```

23         c=c+steady(prefix + j + 1);
24     end
25     total=total+c;
26     if d
27         fprintf( '%d hops:\t%f\n', i, c);
28     end
29     avg=avg+(c*i);
30 end
31
32 loss=steady(len);
33 avg=avg/(1-loss);
34 if d
35     fprintf( 'Loss:\t%f\nTotal:\t%f\nAverage #hops:\t%f\n', loss, total + loss, avg);
36 end
37 end

```

Listing 6: avghops.m

```

1 function [ avg ] = ruavghops( Q, d )
2 %RUAVGHOPS Calculate average number of times a job is forwarded for
   the random unvisited chain
3 %Parameters:
4 %     Q      The matrix representing a markov chain using the
   random unvisited forwarding algorithm
5 %Optional:
6 %     d      Debug mode, default=1, disable debug output=0
7
8     if nargin < 2
9         d=1;
10    end
11
12    steady=ctmcsteadystate(Q);
13
14    len=length(Q);
15
16    avg = 0;
17    avgp = zeros(1, len);
18
19    for i=0:len-2
20        tmpavg = 0;
21        for h=0:i
22            c = prod(i-h+1:i) * (len-1-i) / prod(len-1-
                h:len-1);
23            tmpavg = tmpavg + (c * h);
24            avgp(h+1) = avgp(h+1) + (c*steady(i+1));
25        end
26        avg=avg + steady(i+1) * tmpavg;
27    end
28
29    avgp(len) = steady(len);
30
31    loss=steady(len);
32    avg=avg/(1-loss);
33    if d
34        avgp
35        fprintf( 'Loss:\t%f\nAverage #hops:\t%f\n', loss,
                avg);
36    end
37
38 end

```


Listing 7: ruavghops.m

```

1 function [ pi ] = ctmcsteadystate( Q )
2 %CTMCSTEADYSTATE Steady state distribution of a continious time
   markov chain
3 %Parameters:
4 %      Q      Matrix representing a Markov Chain
5 %Source: http://speed.cis.nctu.edu.tw/~ydlin/course/cn/nsd2009/
   Markov-chain.pdf (slide 10)
6
7      T=Q;
8      len=length(Q);
9      T(:,len)=ones(len, 1);
10     e=zeros(1, len);
11     e(len)=1;
12     pi=e*inv(T);
13 end

```

Listing 8: ctmcsteadystate.m

```

1 function [ avg ] = lumpavghops(Q)
2 %LUMPAVGHOPS Get the average number of times a job is forwarded
   when the state matrix is lumped
3 %Parameters:
4 %      Q      A lumped matrix representation of a markov Chain
5
6      fullsize=length(Q);
7      [Q S]=lump(Q);
8      lumpsize=length(S);
9      nodes=log2(fullsize);
10     hops=zeros(1,nodes+1);
11
12     steady=ctmcsteadystate(Q);
13
14     hops(1)=steady(1); %zero hops
15     hops(nodes+1)=steady(lumpsize); %loss
16     for i=2:lumpsize-1;
17         bits=ceil(log2(S(i)+1));
18         hops(1)=hops(1)+(nodes-bits)/nodes*steady(i);
19
20         for j=bits:-1:1
21             c=0;
22             while c<j && bitand(S(i), 2^(j-c-1))
23                 c=c+1;
24             end
25             hops(c+1)=hops(c+1) + steady(i)/nodes;
26         end
27     end
28
29     %fprintf('Sum:\t%f\n',sum(hops));
30
31     avg=(hops(1:nodes)*[0:nodes-1]')/(1-steady(lumpsize));
32
33 end

```

Listing 9: lumpavghops.m

```

1 function [Ql S] = lump(Q)
2 %LUMP Lump a matrix representing a Markov Chain

```

```

3 %Parameters:
4 %      Q      The matrix that should be lumped
5 %Return:
6 %      Ql      The lumped matrix representation
7 %      S      The states that are used in the lumped matrix
8 %The states of the matrix Q must represent the availability of the
   the servers
9
10      [S R C] = makestates(log2(length(Q)));
11
12      Ql=sparse(length(S), length(S));
13
14      [i j s] = find(Q);
15
16      for x=1:length(i)
17          Ql(R(i(x)),R(j(x)))=Ql(R(i(x)),R(j(x)))+s(x);
18      end
19
20      for x=1:length(S)
21          Ql(x,:)=Ql(x,:)/C(x);
22      end
23
24 end

```

Listing 10: lump.m

```

1 function [r, reindex, coverage] = makestates(rsize)
2 %Generate lumped states
3 %Parameters:
4 %      rsize      Size of the ring (or log2 of the number of
   states of the matrix)
5 %Return:
6 %      r      Vector of the remaining states, ordered
7 %      reindex      Reference index, each old state points to
   the new lumped state
8 %      coverage      How many states the lumped state with the
   same index represents
9
10      powers = 2.^[0:rsize-1];
11
12      function [v] = rotate(a, size)
13          p=2^(size-1);
14          v = a*2 + floor(a/p) - 2*p*floor(a/p);
15      end
16
17      function [r] = makesmallest(a)
18          r=a;
19          for i=1:(rsize-1)
20              a=rotate(a, rsize);
21              if a<r
22                  r=a;
23              end
24          end
25      end
26
27      reindex = [];
28      for i=0:(2^rsize)-1
29          reindex = [reindex makesmallest(i)];
30      end
31
32      function [c] = cover(a, size)
33          c=1;

```

```

34         a=makesmallest(a);
35         b=rotate(a, size);
36         while a ~= b
37             b=rotate(b, size);
38             c=c+1;
39         end
40     end
41
42     function [r] = smallest(a, size)
43         r=a;
44         for i=1:(size-1)
45             a=rotate(a, size);
46             if a<r
47                 r=a;
48             end
49         end
50     end
51
52     function [v] = f(word, bits, place, size)
53         if place > size
54             v = word;
55         elseif bits == 0
56             v = f(word, bits, place+1, size);
57         elseif place + bits > size
58             v = f(word + powers(place), bits-1, place
59                 +1, size);
60         else
61             v = [f(word + powers(place), bits-1, place
62                 +1, size) f(word, bits, place+1, size)
63                 ];
64         end
65     end
66
67     function [r] = makecombs(k, n)
68         leadzeros = ceil(n/k)-1;
69         fullsize = n - leadzeros - 1;
70         r = f(0,k-1,1,fullsize)*2 + 1;
71     end
72
73     r=[0 2^(rsize)-1];
74     for i=1:rsize-1
75         r = [r makecombs(i, rsize)];
76     end
77
78     s=[];
79     for i=r
80         s=[s reindex(i+1)];
81     end
82
83     r=unique(s);
84     reindex = arrayfun(@(x) find(r == x), reindex);
85
86     coverage=[];
87     for w=r
88         coverage=[coverage cover(w, rsize)];
89     end
90 end

```

Listing 11: makestates.m